

Measuring what actually matters

Indian AI teams still pitch accuracy on slide 3. Investors, CXOs and users have moved on. Modern AI metrics pair MODEL quality (precision/recall/hallucination rate) with PRODUCT outcomes (activation, retention, revenue per session). The PM owns the bridge between them.

Last OKR you owned — was the KR a leading or...

WHAT'S INSIDE

- 01** 4-layer AI metric stack
- 02** Leading vs lagging — AI edit
- 03** OKR structure (AI edit)
- 04** Vanity vs value
- 05** Picking an eval — 3 questions

WHY THIS LESSON MATTERS

If your AI metric is "accuracy", you are measuring 2015.

WHY

Real

A real reason this matters in industry today.

SCALE

Today

A real reason this matters in industry today.

SPEED

Now

A real reason this matters in industry today.

IMPACT

Yours

A real reason this matters in industry today.

LEARNING OUTCOMES

By the end of this lesson, you will be able to —

1

Input / Output / Outcome / Impact metric stack.

2

Leading vs lagging metrics for AI.

3

Write 1 AI OKR with 3 KRs.

4

Spot vanity vs value metrics.

5

Pick the right eval for the PRD.

AGENDA

What we'll cover today

Short, sequenced parts. Each one builds on the last.

01

4-layer AI metric stack

Key concept of this lesson.

02

Leading vs lagging — AI edit

You cannot steer with the rear-view mirror

03

OKR structure (AI edit)

Objective + 3 KRs

04

Vanity vs value

Key concept of this lesson.

01

PART 1 OF 4

4-layer AI metric stack

Key concept of this lesson.

4-layer AI metric stack

A core idea you'll use throughout this lesson.

DEFINITION

4-layer AI metric stack

EXAMPLE

Real-world example: applies directly to Indian industry contexts.

Leading vs lagging — AI edit

Leading: eval-set score, golden-set pass, shadow-traffic agreement.

Lagging: monthly retention, revenue per session.

PRD should name BOTH.

Never ship with only lagging metrics.

RULE

Both or neither

Leading drives, lagging confirms.

KEY POINTS

- Leading: eval-set score, golden-set pass, shadow-traffic agreement
- Lagging: monthly retention, revenue per session
- PRD should name BOTH
- Never ship with only lagging metrics

CONCEPT 3 OF 4

OKR structure (AI edit)

Objective = qualitative outcome ("Make support self-serve").

KR1 = model metric ("Answer quality $\geq 4.2 / 5$ ").

KR2 = product metric ("Deflection $\geq 40\%$ ").

KR3 = guardrail ("Hallucination $\leq 2\%$ ").

TEMPLATE

1 + 3

One objective, three KRs.

DEFINITION

OKR structure (AI edit)

Objective + 3 KRs

EXAMPLE

Real-world example: applies directly to Indian industry contexts.

Vanity vs value

A core idea you'll use throughout this lesson.

KEY POINTS

- A core idea you'll use throughout this lesson
- Practical, named, applied to a real Indian use-case.
- Practical, named, applied to a real Indian use-case.

02

PART 2 OF 4

Engage and reflect

Pause, talk, and connect what you just learned to your own work.

✦ THINK • PAIR •
SHARE

ACTIVITY

Write 1 AI OKR

PROMPT

Apply the four concepts above to one real example from your own week.

HOW TO RUN IT

- 1
- 2 Swap with partner, mark vanity vs value.
- 3
- 4 Reveal: instructor confirms the canonical answer.

REFLECT & APPLY

Three questions before you close this lesson

Spend 60 seconds on each. Write the answers down — no scrolling, no shortcuts.

Q1

Which KR in your team is secretly vanity?

Q2

Do users feel your dashboard metric?

Q3

What guardrail does your AI need before Q3?

03

PART 3 OF 4

Knowledge check

Three short questions. One minute each. Honest answers only.

Top layer of metric stack:



Input



Output



Outcome



Impact

Top layer of metric stack:



CORRECT ANSWER

Impact

WHY

Impact is the business-facing top

A leading metric for an AI assistant:

A

Monthly retention

B

Shadow-traffic agreement

C

NPS

D

Revenue

A leading metric for an AI assistant:



CORRECT ANSWER

Shadow-traffic agreement

WHY

Pre-launch, can be measured daily

Good OKR = 1 Objective +:

A

1 KR

B

3 KRs

C

10 KR_s

D

E

Good OKR = 1 Objective +:



CORRECT ANSWER

3 KRs

WHY

Classic "1 + 3" pattern

04

PART 4 OF 4

Apply and close

One thing to remember. One thing to ship this week.

SUMMARY

Five sentences to remember

If you remember nothing else from this lesson, remember these five lines.

- 1 Pick the right tool — never the loudest one.
- 2 Practice the framework on a real example.
- 3 Watch for the three classic pitfalls.
- 4 Document your decision in one sentence.
- 5 Carry the discipline forward to the next lesson.

WORKED EXAMPLE

How this lesson plays out in one real Indian context

THE SCENARIO

A team applies the lesson's main idea to a real workflow — and the results show up within a quarter.

THE TAKEAWAY

"Good AI metrics travel: model → product → business."

WHAT TO NOTICE

1

4-layer AI metric stack

2

Leading vs lagging — AI edit

3

OKR structure (AI edit)

4

Vanity vs value

Map your PRD to OKRs

Attach the OKR block to your Week-10 PRD.

DELIVERABLES

- Write Objective.
- Write 3 KRs across stack.
- Identify guardrail KR.
- Document one decision you made because of this lesson.

WHAT 'GOOD' LOOKS LIKE

Deliverable: OKR block appended to PRD.

ONE THING TO REMEMBER

“

"Good AI metrics travel: model → product → business."

— AI Learning Hub

END OF LESSON 10

What you can now do

Each one of these is now part of your working vocabulary.

- 1 Pick the right tool — never the loudest one.
- 2 Practice the framework on a real example.
- 3 Watch for the three classic pitfalls.
- 4 Document your decision in one sentence.
- 5 Carry the discipline forward to the next lesson.

NEXT → Lesson 3 — AI Roadmaps.