

AGI · scaling · neurosymbolic · your view

Partners, investors and CEOs will ask where this is going. Your answer shapes what you build. You don't have to be right — you have to be thoughtful, informed, and willing to update.

What's your current 2026 AI view in one sentence?

L7

WHAT'S INSIDE

- 01 Scaling hypothesis
- 02 Operational AGI definitions
- 03 Neurosymbolic + post-Transformer
- 04 3 near-term scenarios
- 05 Your personal view

WHY THIS LESSON MATTERS

"AI is uncertain" is not a career plan.

WHY

Real

A real reason this matters in industry today.

SCALE

Today

A real reason this matters in industry today.

SPEED

Now

A real reason this matters in industry today.

IMPACT

Yours

A real reason this matters in industry today.

LEARNING OUTCOMES

By the end of this lesson, you will be able to —

- 1 Scaling hypothesis + critics.
- 2 Operational AGI definitions.
- 3 Neurosymbolic + post-Transformer candidates.
- 4 A personal defensible view.
- 5 Scenario-plan across 3 futures.

AGENDA

What we'll cover today

Short, sequenced parts. Each one builds on the last.

01

Scaling hypothesis

The bet behind GPT-class models

02

Operational AGI definitions

Key concept of this lesson.

03

Neurosymbolic +...

What's brewing

04

3 near-term scenarios

Key concept of this lesson.

01

PART 1 OF 4

Scaling hypothesis

The bet behind GPT-class models

CONCEPT 1 OF 4

Scaling hypothesis

Claim: intelligence emerges predictably with scale.

Evidence: Chinchilla, GPT-4, Claude, Gemini.

Critics: data wall, diminishing returns, cost.

Implication: compute + data = moat.

LAW

Scale → emergence

Holds — so far.

DEFINITION

Scaling hypothesis

The bet behind GPT-class models

EXAMPLE

Real-world example: applies directly to Indian industry contexts.

Operational AGI definitions

A core idea you'll use throughout this lesson.

KEY POINTS

- A core idea you'll use throughout this lesson
- Practical, named, applied to a real Indian use-case.
- Practical, named, applied to a real Indian use-case.

Neurosymbolic + post-Transformer

Pure neural hits provable-reasoning limits.
Neurosymbolic — LLM + solver / planner.
Mamba / SSMs — linear attention, long context.
Diffusion language (SEDD) — early.
Reasoning-as-training (o1) — new scaling axis.
FRONTIER
Watch list
Bet small; read closely.

DEFINITION

Neurosymbolic + post-Transformer

What's brewing

EXAMPLE

Real-world example: applies directly to Indian industry contexts.

3 near-term scenarios

A core idea you'll use throughout this lesson.

KEY POINTS

- A core idea you'll use throughout this lesson
- Practical, named, applied to a real Indian use-case.
- Practical, named, applied to a real Indian use-case.

02

PART 2 OF 4

Engage and reflect

Pause, talk, and connect what you just learned to your own work.

✦ THINK • PAIR •
SHARE

ACTIVITY

Write your 2026 view

PROMPT

Apply the four concepts above to one real example from your own week.

HOW TO RUN IT

1

2

3

4

Reveal: instructor confirms the canonical answer.

REFLECT & APPLY

Three questions before you close this lesson

Spend 60 seconds on each. Write the answers down — no scrolling, no shortcuts.

Q1

What evidence would change your view in 30 days?

Q2

Which thinker do you copy too much?

Q3

Where is your view contrarian?

03

PART 3 OF 4

Knowledge check

Three short questions. One minute each. Honest answers only.

Scaling critic argument:

A

Data wall

B

More compute

C

Fewer params

D

None

Scaling critic argument:



CORRECT ANSWER

Data wall

WHY

Data is finite

Operational AGI test:

A Beauty

B Economic value

C Colour

D None

Operational AGI test:

CORRECT ANSWER



Economic value

WHY

Economic is measurable

Post-Transformer candidate:

A

RNN

B

Mamba

C

BERT

D

LSTM

Post-Transformer candidate:



CORRECT ANSWER

Mamba

WHY

SSMs = Mamba

04

PART 4 OF 4

Apply and close

One thing to remember. One thing to ship this week.

SUMMARY

Five sentences to remember

If you remember nothing else from this lesson, remember these five lines.

- 1 Pick the right tool — never the loudest one.
- 2 Practice the framework on a real example.
- 3 Watch for the three classic pitfalls.
- 4 Document your decision in one sentence.
- 5 Carry the discipline forward to the next lesson.

WORKED EXAMPLE

How this lesson plays out in one real Indian context

THE SCENARIO

Claim: intelligence emerges predictably with scale.

Evidence: Chinchilla, GPT-4, Claude, Gemini.

Critics: data wall, diminishing returns, cost.

Implication: compute + data = moat.

LAW

Scale → emergence

Holds — so far.

THE TAKEAWAY

"Thoughtful beats confident; updated beats thoughtful."

WHAT TO NOTICE

1

Scaling hypothesis

2

Operational AGI definitions

3

Neurosymbolic +...

4

3 near-term scenarios

Your 2026 view

Write a 500-word post.

DELIVERABLES

- Assign scenario weights.
- Name 3 flip beliefs.
- Name 3 signals.
- Document one decision you made because of this lesson.

WHAT 'GOOD' LOOKS LIKE

Deliverable: 500-word post on LMS.

ONE THING TO REMEMBER

“

"Thoughtful beats confident; updated beats thoughtful."

— AI Learning Hub

END OF LESSON 11

What you can now do

Each one of these is now part of your working vocabulary.

- 1 Pick the right tool — never the loudest one.
- 2 Practice the framework on a real example.
- 3 Watch for the three classic pitfalls.
- 4 Document your decision in one sentence.
- 5 Carry the discipline forward to the next lesson.

NEXT → Projects & Assessment.